


Range Searching


- Data structure for a set of objects (points, rectangles, polygons) for efficient range queries.



- Depends on type of objects and queries. Consider basic data structures with broad applicability.
- Time-Space tradeoff: the more we preprocess and store, the faster we can solve a query.
- Consider data structures with (nearly) linear space.

Orthogonal Range Searching


- Fix a n -point set P . It has 2^n subsets. How many are possible answers to geometric range queries?



- Efficiency comes from the fact that only a small fraction of subsets can be formed.
- Orthogonal range searching deals with point sets and axis-aligned rectangle queries.
- These generalize 1-dimensional sorting and searching, and the data structures are based on compositions of 1-dim structures.

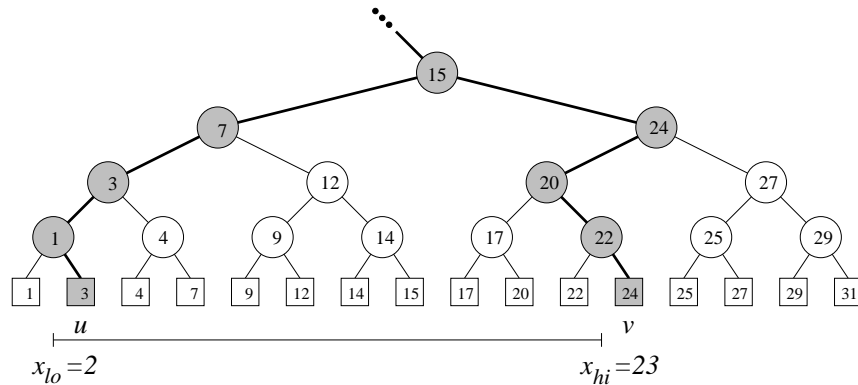
1-Dimensional Search

- Points in 1D $P = \{p_1, p_2, \dots, p_n\}$.
- Queries are intervals.



- If the range contains k points, we want to solve the problem in $O(\log n + k)$ time.
- Does hashing work? Why not?
- A sorted array achieves this bound. But it doesn't extend to higher dimensions.
- Instead, we use a balanced binary tree.

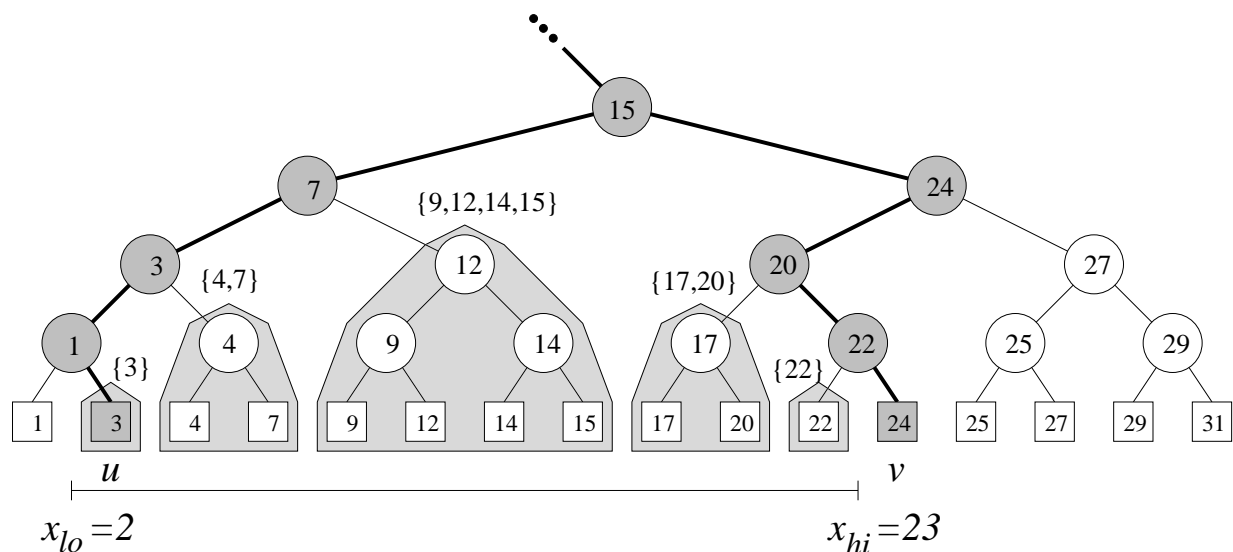
Tree Search



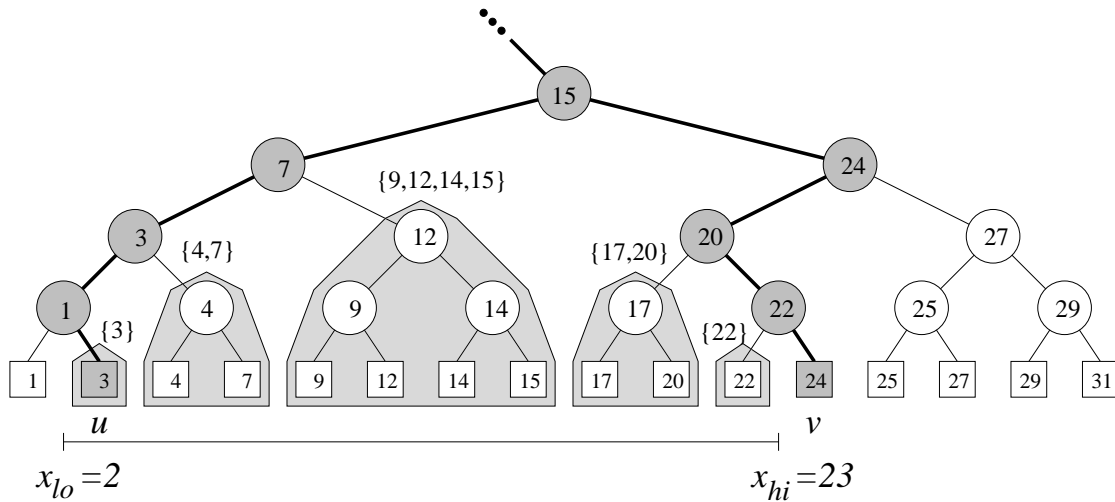
- Build a balanced binary tree on the sorted list of points (keys).
- Leaves correspond to points; internal nodes are branching nodes.
- Given an interval $[x_{lo}, x_{hi}]$, search down the tree for x_{lo} and x_{hi} .
- All leaves between the two form the answer.
- Tree search takes $2 \log n$, and reporting the points in the answer set takes $O(k)$ time; assume leaves are linked together.

Canonical Subsets

- S_1, S_2, \dots, S_k are **canonical subsets**, $S_i \subseteq P$, if the answer to **any** range query can be written as the disjoint union of some S_i 's.
- The canonical subsets may overlap.
- Key is to determine correct S_i 's, and given a query, efficiently determine the appropriate ones to use.
- In 1D, a canonical subset for each node of the tree: S_v is the set of points at the leaves of the subtree with root v .




1D Range Query



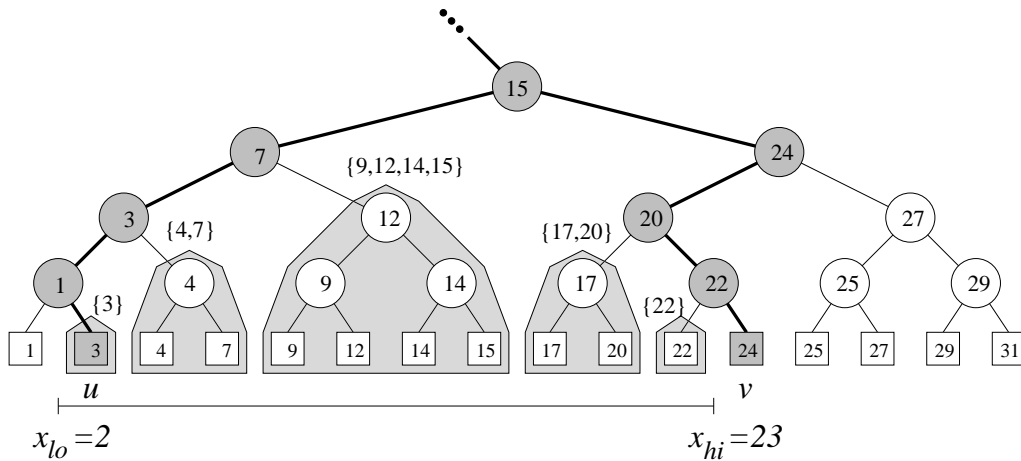
- Given query $[x_{lo}, x_{hi}]$, search down the tree for **leftmost leaf** $u \geq x_{lo}$, and **leftmost leaf** $v \geq x_{hi}$.
- All leaves between u and v are in the range.
- If $u = x_{lo}$ or $v = x_{hi}$, include that leaf's canonical set (singleton) into the range.
- The remainder range determined by **maximal** subtree lying in the range $[u, v]$.

Query Processing

- Let z be the last node common to search paths from root to u, v .
- Follow the left path from z to u . When path goes left, add the canonical subset of right child. (Nodes 7, 3, 1 in Fig.)
- Follow the right path from z to v . When path goes right, add the canonical subset of left child. (Nodes 20, 22 in Fig.)




Analysis




- Since search paths have $O(\log n)$ nodes, there are $O(\log n)$ canonical subsets, which are found in $O(\log n)$ time.
- To list the sets, traverse those subtrees in linear time, for additional $O(k)$ time.
- If only **count** is needed, storing sizes of canonical sets at nodes suffices.
- Data structure uses $O(n)$ space, and answers range queries in $O(\log n)$ time.

Multi-Dimensional Data




- Range searching in higher dimensions?
- kD -trees [Jon Bentley 1975]. Stands for **k -dimensional trees**.
- Simple, general, and arbitrary dimensional. Asymptotic search complexity not very good.
- Extends 1D tree, but alternates using x - y -coordinates to split. In k -dimensions, cycle through the dimensions.




- A binary tree. Each node has two values: split dimension, and split value.
- If split along x , at coordinate s , then left child has points with x -coordinate $\leq s$; right child has remaining points. Same for y .
- When $O(1)$ points remain, put them in a leaf node.
- Data points at leaves only; internal nodes for branching and splitting.

Splitting




- To get balanced trees, use the **median** coordinate for splitting—median itself can be put in either half.
- With median splitting, the height of the tree guaranteed to be $O(\log n)$.
- Either cycle through the splitting dimensions, or make data-dependent choices. E.g. select dimension with max spread.

Space Partitioning View



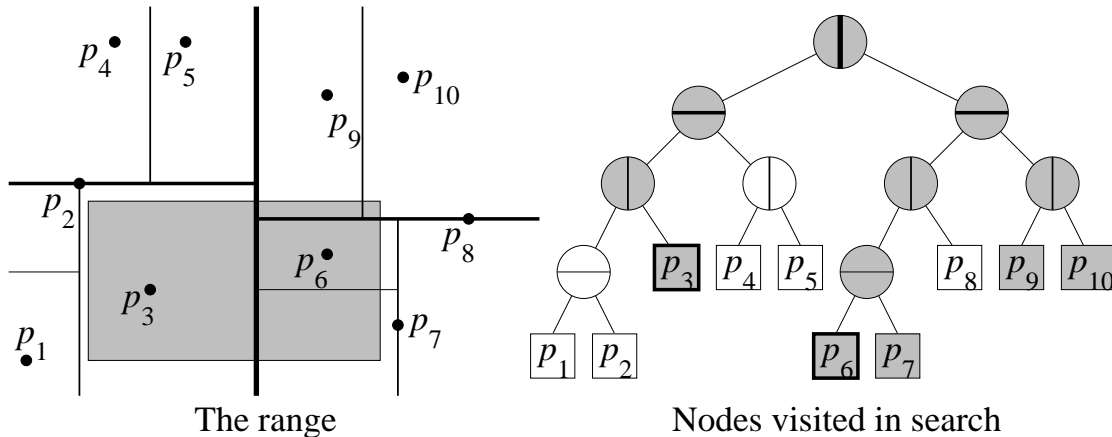
- kD -tree induces a space subdivision—each node introduces a x - or y -aligned cut.
- Points lying on two sides of the cut are passed to two children nodes.
- The subdivision consists of rectangular regions, called **cells** (possibly unbounded).
- Root corresponds to entire space; each child inherits one of the halfspaces, so on.
- Leaves correspond to the terminal cells.
- Special case of a general partition BSP.

Construction




- Can be built in $O(n \log n)$ time recursively.
- Presort points by x and y -coordinates, and cross-link these two sorted lists.
- Find the x -median, say, by scanning the x list. Split the list into two. Use the cross-links to split the y -list in $O(n)$ time.
- Now two subproblems, each of size $n/2$, and with their own sorted lists. Recurse.
- Recurrence $T(n) = 2T(n/2) + n$, which solves to $T(n) = O(n \log n)$.

Searching kD -Trees




- Suppose query rectangle is R . Start at root node.
- Suppose current splitting line is vertical (analogous for horizontal). Let v, w be left and right children nodes.
- If v a leaf, report $cell(v) \cap R$;
if $cell(v) \subseteq R$, report all points of $cell(v)$;
if $cell(v) \cap R = \emptyset$, skip;
otherwise, search subtree of v recursively.
- Do the same for w .
- Procedure obviously correct. What is the time complexity?

Search Complexity



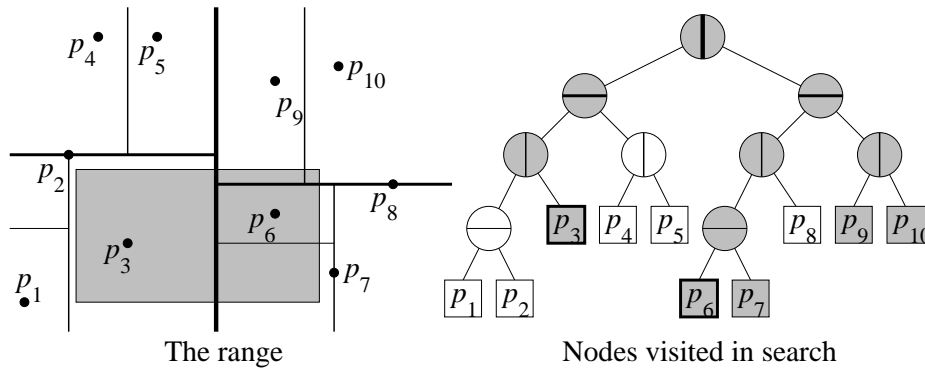
- When $cell(v) \subseteq R$, complexity is linear in output size.
- It suffices to bound the number of nodes v visited for which the boundaries of $cell(v)$ and R intersect.
- If $cell(v)$ outside R , we don't search it; if $cell(v)$ inside R , we enumerate all points in region of v ; a recursive call is made only if $cell(v)$ partially overlaps R ; the kD -tree height is $O(\log n)$.
- Let ℓ be the line defining one side of R .
- We prove a bound on the number of cells that intersect ℓ ; this is more than what is needed; multiply by 4 for total bound.

Search Complexity



- How many cells can a line intersect?
- Since splitting dimensions alternate, the key idea is to consider two levels of the tree at a time.
- Suppose the first cut is vertical, and second horizontal. We have 4 cells, each with $n/4$ points.
- A line intersects exactly two cells; the others cells will be either outside or entirely inside R .
- The recurrence is

$$Q(n) = \begin{cases} 1 & \text{if } n = 1, \\ 2Q(n/4) + 2 & \text{otherwise.} \end{cases}$$



- **The recurrence $Q(n) = 2Q(n/4) + 2$ solves to**

$$Q(n) = O(\sqrt{n})$$

- kD -Tree is an $O(n)$ space data structure that solves 2D range query in worst-case time $O(\sqrt{n} + m)$, where m is the output size.

d -Dim Search Complexity

- What's the complexity in higher dimensions?
- Try 3D, and then generalize.
- The recurrence is


$$Q(n) = 2^{d-1}Q(n/2^d) + 1$$

- It solves to

$$Q(n) = O(n^{1-1/d})$$

- kD -Tree is an $O(dn)$ space data structure that solves d -dim range query in worst-case time $O(n^{1-1/d} + m)$, where m is the output size.

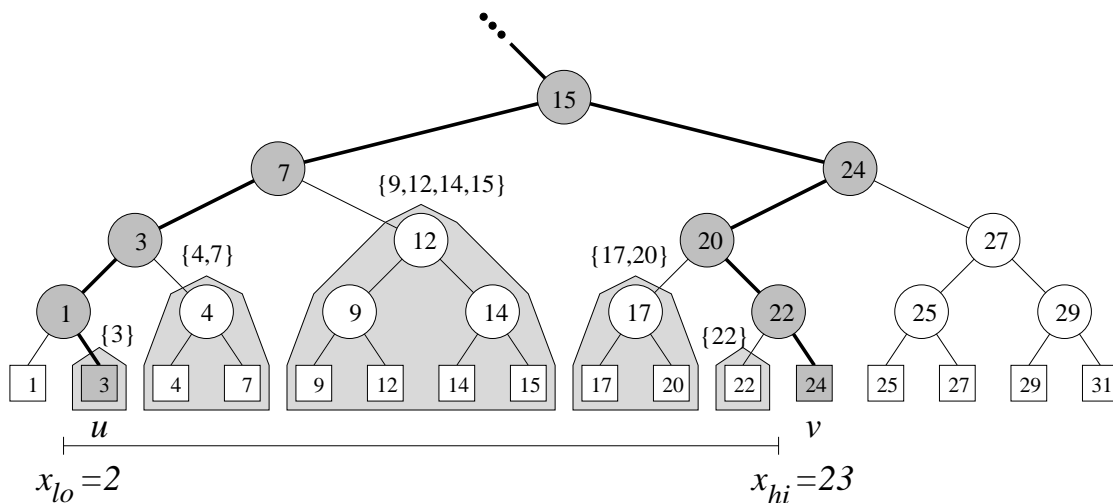
Orthogonal Range Trees



- Generalize 1D search trees to dimension d .
- Each search recursively decomposes into multiple lower dimensional searches.
- Search complexity is $O((\log n)^d + k)$, where k is the answer size.
- Space & time complexity $O(n(\log n)^{d-1})$.
- Fractional cascading eliminates one $\log n$ factor from search time.
- We focus on 2D, but ideas readily extend.

2D Range Trees

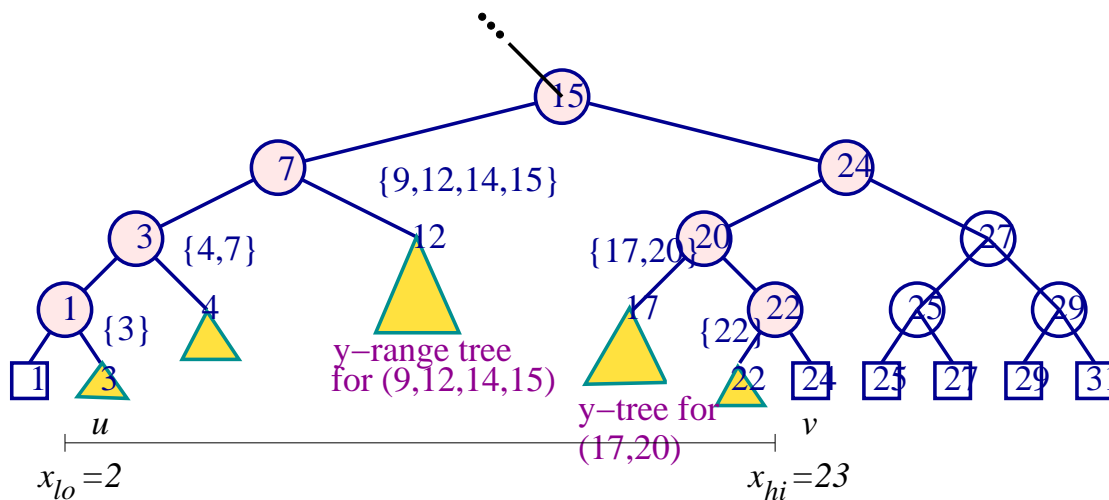
- Suppose $P = \{p_1, p_2, \dots, p_n\}$ set of points in the plane.
- The generic query is $R = [x_{lo}, x_{hi}] \times [y_{lo}, y_{hi}]$.
- We first ignore the y -coordinates, and build a 1D x -range tree on P .



- The set of points that fall in $[x_{lo}, x_{hi}]$ belong to $O(\log n)$ canonical sets.
- This is a superset of the final answer. It can be significantly bigger than $|R \cap P|$, so we can't afford to look at each point in these canonical sets.


Level 2 Trees

- Key idea is to collect points of each canonical set, and build a **y-range tree** on them.
- E.g., the canonical set $\{9, 12, 14, 15\}$ is organized into a 1D range tree using those points' y -coordinates.



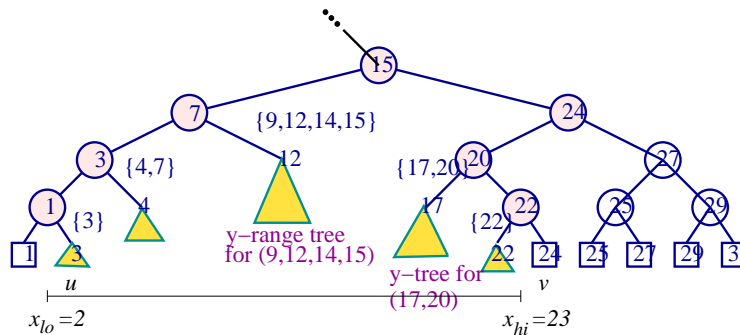
- We search each of the $O(\log n)$ canonical sets that include points for x -range $[x_{lo}, x_{hi}]$ using their y -range trees for range $[y_{lo}, y_{hi}]$.
- The y -range searches list out the points in $R \cap P$. (No duplicates.)

Canonical Sets



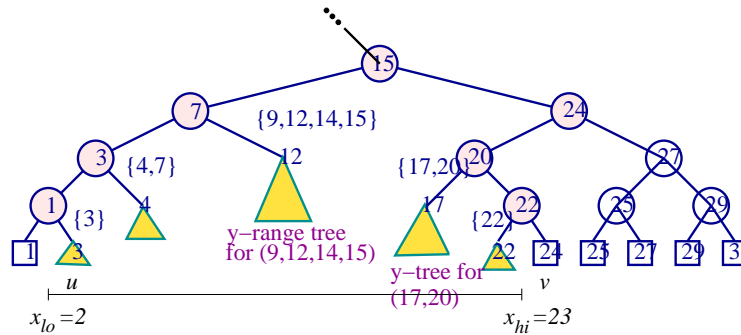
Analysis

- Time complexity for 2D is $O((\log n)^2)$.
 1. $O(\log n)$ canonical sets for x -range.
 2. Each set's y -range query takes $O(\log n)$ time.



- Space complexity is $O(n \log n)$.
 1. What is the total size of all canonical sets in x -tree?
 2. Number of nodes \equiv number of leaves.
 3. One set of size n . Two of size $n/2$, etc.
 4. Total is $O(n \log n)$.
 5. Each canonical set of size m requires $O(m)$ space for the y -range tree.
 6. So, overall space is $O(n \log n)$.

Construction



- The x -tree can be built in $O(n \log n)$ time.
- Naively, since total size of all y -trees is $O(n \log n)$, it will take $O(n(\log n)^2)$ time to build them.
- By building them bottom-up, we can avoid sorting cost at each node.
- Once y -trees for the children nodes are built, we can merge their y -lists to get the parent's y -list in linear time.
- The cost of building the 1D range tree is linear after sorting.
- Thus, total time is linear in $O(n \log n)$, the total sizes of all y -trees.

d -Dim Range Trees


- The multi-level range tree idea extends naturally to any dimension d .
- Build the x -tree on first coordinate.
- At each node v of this tree, build the $(d - 1)$ -dimensional range tree for canonical set of v on the remaining $d - 1$ dimensions.
- Search complexity grows by one $\log n$ factor for each dimension—each dimensional increases the number of canonical sets by $\log n$ factor.
- So, search cost is $O((\log n)^d)$.
- Space and time complexity is $O(n(\log n)^{d-1})$.

Fractional Cascading

- A technique that improves the range tree search time by log factor. 2D search can be done in $O(\log n)$ time.
- Basic idea: Range tree first finds the set of points lying in $[x_{lo}, x_{hi}]$ as union of $O(\log n)$ canonical sets.
- Next, each canonical set is searched using the y -tree for range $[y_{lo}, y_{hi}]$. We locate y_{lo} ; then read off points until y_{hi} reached.
- Since each set is searched for the **same** key, y_{lo} , we can improve the search to $O(1)$ per set.
- In effect, we do the first search in $O(\log n)$ time, but then use that information to search other structures more efficiently.
- The key is to place smart hooks linking the search structures for the canonical sets.


Basic Idea

- To understand the basic idea, consider a simple example.
- We have two sets of numbers, A_1, A_2 , both sorted.
- Given a range $[x, x']$, want to report all keys in A_1, A_2 that lie in the range.
- Straightforward method takes $2 \log n + k$, if k is the answer size; separate binary searches in A_1, A_2 to locate x .
- For example, range $[20, 65]$.



Fractional Cascading Idea

- Suppose $A_2 \subset A_1$. Add pointers from A_1 to A_2 .
- If $A_1[i] = y_i$, store ptr to entry in A_2 with smallest key $\geq y_i$. (Nil if undefined.)





- Suppose we want keys in range $[y, y']$.
- Search A_1 for y , and walk until past y' . Time $O(\log n + k_1)$.
- If A_1 search for y ended at $A_1[i]$, use its pointer to start search in A_2 . This takes $O(1 + k_2)$ time.
- Example $[20, 65]$.

FC in Range Trees

- **Key observation:** canonical subsets $S(\ell(v))$ and $S(r(v))$ are subsets of $S(v)$.
- The x -tree is same as before. But instead of building y -trees for canonical subsets, we store them as sorted arrays, by y -coordinate.
- Each entry in $A(v)$ stores two pointers, into arrays $A(\ell(v))$ and $A(r(v))$.
- If $A(v)[i]$ stores point p , then ptr into $A(\ell(v))$ is to entry with smallest y -coordinate $\geq y(p)$. Same for $(r(v))$.


Range Tree FC

- Only some pointers shown to avoid clutter.



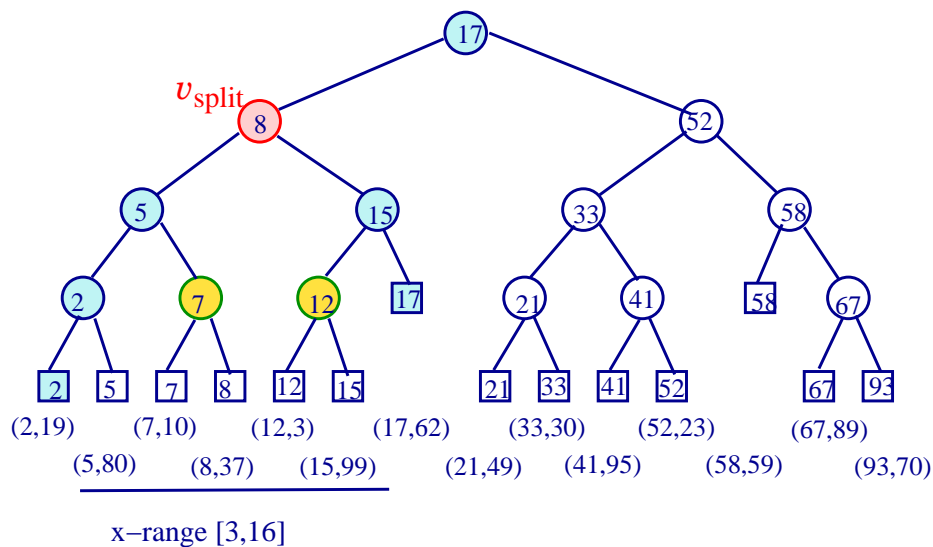
FC Search

- Consider range $R = [x, x'] \times [y, y']$.
- Search for x, x' in the main x -tree.
- Let v_{split} be the node where the two search paths diverge.
- The $O(\log n)$ canonical subsets correspond to nodes that lie below v_{split} , and are the right (left) child of a node on search path to x (resp. x') where the path goes left (resp. right).



FC Search


- At v_{split} , do binary search to locate y in $A(v_{split})$. $O(\log n)$ time.
- As we search down the x -tree for x, x' , keep track of the entries in the associated arrays for smallest keys $\geq y$, at $O(1)$ cost per node.



- Let $A(v)$ be one of the $O(\log n)$ canonical nodes that is to be searched for $[y, y']$ range.
- We just need to find the smallest entry in $A(v) \geq y$.

FC Search

- We can find this in $O(1)$ time because $\text{parent}(v)$ is on the search path, and we know smallest entry $\geq y$ in $A(\text{parent}(v))$, and have a pointer from that to v 's array.



- So we can output all points in $A(v)$ that lie in range $[y, y']$ in time $O(1 + k_v)$, where k_v is the answer size.
- For 2D range search, the final time complexity is $O(\log n + k)$, and space $O(n \log n)$.
- d -dim range search takes $O((\log n)^{d-1} + k)$ time with fractional cascading.